

An Explainable Physics-Guided Digital Twin for Multi-Horizon Glucose Forecasting in Type 1 Diabetes

Samyak Jain, Dhriti Goswami, Yugnanda Malhotra

Abstract—Background and objective: Glucose dynamics in type 1 diabetes (T1D) are nonlinear, delayed, and patient-specific, influenced by meals and insulin. Clinically meaningful forecasting requires valid prediction horizons and comparison with persistence baselines. This study proposes an explainable, physics-guided digital twin integrating deep sequence learning with mechanistic physiological modeling for personalized, multi-horizon glucose forecasting.

Methods: A gap-aware pipeline was developed using the OhioT1DM dataset, ensuring accurate horizon-specific glucose targets. The model combines a compact Transformer encoder with a differentiable Bergman Minimal Model, extended with subcutaneous insulin and gut absorption states, patient-specific parameter estimation, and a non-crossing quantile head for uncertainty. Twelve subjects and 26,498 test windows were evaluated under personalized and subject-disjoint leave-one-subject-out (LOSO) protocols. Performance was assessed across 30–120-minute horizons using RMSE, skill over persistence, Clarke/Parkes error distributions, hypoglycemia-alarm sensitivity, and feature attribution.

Results: The model achieved RMSE values of 18.84, 30.52, 38.60, and 44.46 mg/dL at 30, 60, 90, and 120 minutes, respectively, with a 19.4–21.2% improvement over persistence. LOSO validation confirmed the personalization while the physics constraints preserved accuracy and enabled estimation of patient-specific insulin sensitivity. The lower-quantile alarm further improved hypoglycemia sensitivity.

Conclusion: Physiological constraints with AI-based forecasting support personalized diabetes monitoring. The proposed framework demonstrates potential for clinical translation through stable insulin sensitivity estimation and improved hypoglycemia prediction, while supporting interpretable, physics-guided forecasting.

Index Terms—continuous glucose monitoring, digital twin, type 1 diabetes, physics-informed learning, Transformer, uncertainty quantification, explainable artificial intelligence

I. INTRODUCTION

A. Background

Type 1 diabetes (T1D) is a disease involving the immune-mediated destruction pancreatic β -cells and thus individuals with type 1 diabetes require lifelong exogenous insulin replacement and glucose monitoring, disease-specific education, diet, and lifestyle modification remain cornerstones of its management alongside insulin therapy. Glucose-insulin dynamics, however, are governed by nonlinear, patient-specific physiology, as described mechanistically through the Bergman Minimal Model [1]. Despite advancements in continuous glucose monitoring (CGM) hardware, it remains very difficult to build a system that converts raw diabetes data into personalized, actionable medical recommendations [2].

This has driven the development of machine-learning approaches for CGM forecasting progressing from black-box sequence models, including recurrent networks to attention-based Transformer architectures [3], [4], which can achieve strong short-horizon accuracy but function as black boxes offering clinicians little interpretability [3], [5]. This limitation motivated a shift toward increasingly physiology-aware approaches that fold a physics-consistency penalty derived from mechanistic models, following the Physics-Informed Neural Network (PINN) paradigm [6], [7], [8], [9], [10]. However, validation has typically remained simulator-only or narrowly dataset-specific [7], [8]. More recent systems introduced patient-specific digital twins and prospective evaluation [11], [12], [13]. Yet translating predictions into safe, actionable recommendations remains challenging, with reinforcement-learning approaches still considered insufficiently clinically safe [14].

The field has therefore advanced along several parallel tracks; interpretable risk prediction, physiology-constrained forecasting, digital-twin construction and validation, and increasingly capable sequence-modelling backbones that remain largely disconnected from one another.

B. Literature Review

Systematic reviews of Type 1 diabetes digital twins note that existing systems are not yet comprehensive, multi-scale virtual replicas, and most model glucose-insulin metabolism in isolation rather than the full patient [2], motivating a more integrated, personalized digital twin.

A natural first step was showing CGM-based risk prediction could be both accurate and interpretable. Explainable short-horizon hypoglycemia/hyperglycemia prediction models added SHAP-based feature attributions to opaque risk classifiers, demonstrating strong, interpretable patient-level forecasting [3]. This explained *why* a risk was predicted but not what to do about it, motivating subsequent work using SHAP attributions to guide bounded, one-step behavioural recommendations on OhioT1DM [6]. However, the approach remained exploratory and was not considered clinically safe, shifting subsequent focus toward postprandial glucose response. Explainable postprandial glucose modeling attributed meal-response forecasts to specific inputs, achieving personalized, transparent prediction [7], though remaining purely data-driven with no physiological constraint.

Next, physiology was embedded directly into the learning objective. Early PINN-based glucose predictors added a

physics-consistency penalty to sequence learning, improving short-term in-silico prediction while preserving mechanistic plausibility, though validation stayed simulator-only [9]. Bergman-constrained PINN models tied the network explicitly to the Bergman Minimal Model, improving physiological plausibility on real T1DM data, though evaluation remained narrow and dataset-specific [10]. Hybrid physiology and deep learning architectures embedded glucose–insulin dynamics into both architecture and training, achieving reliable predictions under limited real-subject data, but remained forecasting models rather than virtual patients [11]. Physics-guided sequence models addressed recurrent black-box models’ inability to distinguish meal and insulin effects via physics-LSTM hybrids, bridging sequence modeling and physiology-guided forecasting, still without a full virtual-patient representation [12].

By integrating physiology-based models with LSTM learning, it produced patient-specific virtual subjects for in-silico therapy testing, moving from forecasting toward simulation, though remaining a research framework rather than a clinically validated system [13]. An open-source digital-twin methodology addressed this by identifying a personalized glucose–insulin model directly from routine CGM, insulin, and meal data, showing non-invasive, CGM-only twinning suffices for personalized counterfactual simulation, though use remained mainly retrospective [14]. Simulation-to-reality data augmentation then used simulated patients as synthetic data generators to offset limited T1D data availability, improving glucose prediction model training without yet proving twins support decision-making [15]. Twin fidelity was addressed by validating physiologically constrained digital twins against real exercise data from 394 cases – strong evidence simulated twins match observed clinical dynamics, though still simulation-centric [16]. This gap was partly closed by a randomized clinical trial of a twin-enhanced decision-support system, showing improved time-in-range and reduced hypoglycaemia within a bolus-dosing-centered system [17], further extended to full adaptive biobehavioural co-adaptation for automated insulin delivery (AID) in a six-month randomized trial, where each participant’s AID data generated a cloud-based digital twin biweekly re-optimizing therapy parameters, overcoming the typical time-in-range plateau – though optimizing therapy periodically rather than forecasting glucose continuously in real time [18].

In parallel, attention-based Transformer forecasting on OhioT1DM outperformed classical machine-learning baselines (XGBoost, 1D-CNN) for short-horizon prediction but remained a plain sequence model with limited interpretability [19], extended to interpretable multi-horizon forecasting with attention over temporal and static covariates, though without benchmarking against an explicit recurrent-hybrid baseline [20]. A hybrid Transformer-LSTM architecture addressed pure attention models’ under-modeling of short-term fluctuations, improving short/medium-horizon accuracy, though horizon coverage stopped short of 120 minutes [21]. Long-horizon hybrid Transformer-LSTM models pushed prediction to 120 minutes with richer clinical inputs and strong safety-zone performance, but remained purely data-driven [22]. Metadata-aware personalization improved individualized forecasting via

physiological signals, demographics, and device metadata, but required centralized pooling of sensitive data [23], addressed by personalized federated learning training models across sites without centralizing data, focused on high-risk glycaemic-excursion regions, though not addressing on-device efficiency or interpretability [24]. Practical deployability was addressed via lightweight, edge-deployable Transformer models, still lacking native uncertainty quantification or interpretability [25]. A Temporal Fusion Transformer with built-in attention-based feature/temporal importance and calibrated predictive uncertainty addressed this directly, outperforming prior baselines by $\geq 13\%$ RMSE, though attention weights weren’t validated against physiological ground truth [26]. Self-supervised Transformer pretraining produced reusable glucose representations supporting multiple downstream tasks (screening, dietary prediction), but lacked grounding in verified clinical guidelines [27], addressed by a retrieval-augmented generation framework for diabetes education, retrieving verified medical content before generation to reduce hallucination and improve response accuracy, though applied to conversational education rather than numerical forecasting [28]. Most recently, retrieval-augmented forecasting applied this contextualization directly to CGM prediction, drawing on similar historical episodes rather than raw sequences alone, but still requiring mechanistic grounding, SHAP-style explanation, and prospective clinical validation [29].

Table I situates this study among OhioT1DM forecasting approaches, grouped into physics-informed, non-physics learned/statistical, physics-informed LSTM, and physics-informed hybrid architectures, with the proposed method in the latter category.

C. Objectives

- To jointly train a data-driven Transformer forecasting module and a Bergman Minimal Model-constrained PINN module such that multi-horizon (30, 60, 120 minute) predictions remain consistent with glucose–insulin governing equations.
- To continuously estimate patient-specific physiological parameters (e.g., insulin sensitivity) from routinely collected CGM, insulin, and meal data, without reliance on invasive testing.
- To generate SHAP-based feature attributions for each forecast, corroborated where possible against complementary explanation approaches, and to incorporate retrieval-augmented, guideline-grounded contextual reasoning to support clinician and patient trust.
- To validate the framework on real patient data rather than simulator-only cohorts.

The paper is structured as follows. Section 1 discusses the detailed introduction, literature covering recent advancements in digital twins. Section 2 details the proposed research methodology including data acquisition and model development. Section 3 discusses about the experimental results, model evaluation parameters and real time validation outcomes. Section 4 covers the future prospects of the research including potential expansion for detecting other. Finally, the

TABLE I
LANDSCAPE COMPARISON OF OHIO T1DM BLOOD-GLUCOSE FORECASTING STUDIES BY ARCHITECTURE FAMILY (30- AND 60-MINUTE HORIZONS)

Study	Model family	Cohort	RMSE@30	MAE@30	RMSE@60	MAE@60
<i>Physics-informed only, no learned sequence backbone</i>						
[9]	PINN, short-term glucose	simulator	<i>not numerically comparable (simulator-only validation)</i>			
[10]	PINN under Bergman constraints	non-OhioT1DM	<i>not numerically comparable (different cohort/dataset)</i>			
<i>Non-physics learned/statistical sequence models - 2020 challenge cohort (6 subjects, official split)</i>						
[63]	CNN/LSTM	2020	17.45	11.22	33.67	23.25
[62]	N-BEATS + BiLSTM	2020	18.22	12.83	31.66	23.60
[61]	LSTM (non-personalised)	2020	18.23	14.37	31.10	25.75
[65]	GAN (GRU + 1D-CNN)	2020	18.34	13.37	32.31	24.20
[60]	Shallow NN + error imputation	2020	18.63	10.08	32.27	17.69
[66]	Stacked regression + activity	2020	18.99	13.73	33.39	25.04
[64]	Multi-scale LSTM	2020	19.05	13.50	32.03	23.83
[67]	Multi-lag stacking	2020	19.21	13.93	33.65	25.31
[68]	Latent-variable statistical	2020	19.37	13.76	32.59	24.64
[69]	Multi-class LSTM (+2018 data)	2020	19.40	13.90	33.40	25.00
[70]	Genetic programming	2020	19.60	14.25	32.04	23.58
[71]	Multitask CRNN	2020	19.79	13.62	33.73	24.54
[69]	Multi-class LSTM	2020	19.80	14.40	34.00	25.80
[72]	Online ARMA + residual net	2020	20.03	14.52	34.89	24.61
[73]	Personalised interpretable LSTM	2020	20.20	14.74	34.19	25.98
[71]	Single-task CRNN (ablation)	2020	20.67	14.28	34.40	24.67
[74]	Seq2Seq BiLSTM	2020	21.80	15.00	35.00	25.00
<i>Non-physics learned/statistical sequence models - post-challenge, all-12-subject protocol[†]</i>						
[75]	GATv2 + GRU (GARNN)	all-12	18.97	13.34	—	—
[75]	NHiTS (baseline)	all-12	20.14	14.07	—	—
[75]	N-BEATS (baseline)	all-12	20.15	14.11	—	—
[75]	IMV-LSTM (baseline)	all-12	20.15	14.00	—	—
[75]	RETAIN (baseline)	all-12	20.30	14.41	—	—
[75]	Linear regression (baseline)	all-12	22.19	15.92	—	—
<i>Physics-informed + LSTM hybrid</i>						
[59]	Physiology-informed encoder + LSTM (NPE+LSTM)	2018	17.80	—	—	—
[12]	Physics-LSTM hybrid	—	<i>not numerically comparable (no OhioT1DM value reported)</i>			
<i>Physics-informed + other-architecture hybrid</i>						
[58]	Physiological model + ANN	2018	19.33	—	31.72	—
[11]	Physics-informed glucose-insulin NN	—	<i>not numerically comparable (dataset not confirmed as OhioT1DM)</i>			
[16]	Physiologically-constrained NN (10-compartment)	T1DEXI	<i>not numerically comparable (different dataset)</i>			
—	Current Work: Physics-guided PINN (Bergman) + Transformer	2020	18.74	13.24	31.16	22.55
—	<i>Persistence baseline (ours)</i>	2020	24.22	17.45	40.11	29.51

[†] Cohorts/protocols vary by block: “2020” and “2018” denote the respective 6-subject official-split cohorts, while “all-12” denotes evaluation across all 12 subjects under each study’s protocol. Values are not directly comparable across cohorts. The 90- and 120-minute horizons are omitted because no protocol-matched published OhioT1DM values exist. Table IV reports our own 30-120 min results against our own persistence baseline.

appendix provides supplementary information and supporting details.

II. METHODOLOGY

A. Audit and Design Requirements

The rebuilt system began with an audit of an earlier pipeline. The main faults were: physics disabled in every checkpoint-producing script; an algebraic rather than dynamic remote-insulin state; gap-spanning windows with mislabeled horizons; forward-filled CGM values used as targets; randomly split,

strongly overlapping windows; defective Clarke-grid logic; a time-reversed insulin-on-board kernel; and no persistence baseline. The earlier reported 30-minute RMSE of 30.4 mg/dL was worse than the published persistence value of approximately 22.5 mg/dL.

Each failure became a design constraint or test. The rebuilt implementation contains 234 targeted tests, begins evaluation with persistence, separates the unit of analysis (subject) from the number of overlapping windows, and regenerates all tables from stored predictions.

B. Dataset and Temporal Integrity

OhioT1DM contains 12 subjects in the 2018 and 2020 cohorts and 166,443 CGM observations sampled nominally every five minutes, together with basal and bolus insulin, carbohydrate, activity, sleep, and wearable channels [30], [31]. The official split follows the same subjects forward in time and therefore measures personalized forecasting, not cross-subject generalization. The 2020 cohort's first test hour is excluded according to its challenge protocol; the 2018 cohort is not modified in this way.

All streams are aligned to an exact five-minute grid without dropping rows. Missing CGM runs of at most ten minutes may be interpolated for input construction, but interpolated values are never accepted as targets. A sequence is emitted only if its 24-step (two-hour) input satisfies the coverage rule and every target is a real observation exactly $h \in \{30, 60, 90, 120\}$ minutes after the anchor. The scaler is fit only on training-fold slots, with each physical slot counted once rather than once per overlapping window.

From 187,780 candidate windows, 141,100 (75.1%) survive: 36,395 are rejected for incomplete inputs and 10,285 for unavailable real-valued targets. Both evaluation protocols use the same 26,498 test windows. Because consecutive windows share 23 of 24 input steps, the effective information content is approximately

$$n_{\text{eff}} \approx \frac{135,000 \text{ slots}}{48 \text{ steps}} \approx 2,800, \quad (1)$$

not the raw window count.

C. System Overview

Fig. 1 outlines the full pipeline from raw recordings to reported outcomes. Data collected from all twelve OhioT1DM subjects, comprising continuous glucose readings, insulin delivery, and meal logs, are first aligned to a uniform five-minute time grid; genuine gaps in recording are preserved rather than interpolated over, so that no synthetic values enter the pipeline. From this aligned record, two parallel processes follow. The grid-aligned insulin and meal signals are passed through a simplified physiological model of glucose regulation, based on the Bergman minimal model together with subcutaneous insulin and gut absorption compartments, which yields time-varying estimates of insulin on board, carbohydrate on board, remote insulin action, and the rate of glucose appearance. These physiological estimates are then combined with the glucose signal itself, insulin and meal timing, and time-of-day information to form a single 35-feature representation at each time step.

This feature representation is grouped into two-hour input windows, and windows are retained only when both the input span and the corresponding forecast target are fully intact, so that no window silently spans a genuine data gap. The resulting windows are evaluated under two complementary protocols: an official, personalized setting in which each subject's earlier data are available for training, and a twelve-fold leave-one-subject-out setting in which the model is evaluated on subjects it has never encountered. Each window is then processed by a

compact Transformer encoder, which produces two outputs in parallel: an estimate of the subject's individual physiological parameters, and a probabilistic glucose forecast expressed as a median trajectory bounded by non-crossing lower and upper quantiles.

These two outputs are not independent. The estimated physiological parameters are used to re-simulate the same mechanistic equations introduced earlier, and the resulting trajectory is compared against the network's own forecast at multiple points across the two-hour horizon; the discrepancy between the two constitutes a physics-consistency signal that is folded back into training, alongside a regularization term that keeps the estimated parameters within physiologically plausible bounds. In this way, the mechanistic model contributes to the pipeline twice: once as an input feature, and once as a differentiable constraint on the learned forecast. Finally, forecasts from both evaluation protocols are scored on a per-subject basis and summarized across the cohort, using both conventional accuracy metrics and clinically oriented measures of hypoglycemia detection.

D. Mechanistic Patient Model

The physiological core extends the Bergman minimal model and its insulin-sensitivity index [1], [32]:

$$\dot{G} = -(p_1 + X)G + p_1G_b + \frac{R_a(t)}{V_G}, \quad (2)$$

$$\dot{X} = -p_2X + p_3(I - I_b), \quad (3)$$

$$\dot{I} = -nI + \frac{1000}{V_I} \cdot \frac{S_2}{t_{\max,I}}. \quad (4)$$

Here G is glucose, I plasma insulin, X remote insulin action, and subscript b denotes basal state. All rates are per minute. Basal subtraction in (3) is essential: at $I = I_b$, $X^* = 0$ and $G = G_b$ is an equilibrium. Driving the equation with total insulin instead would create a spurious downward glucose attractor.

The patient-specific insulin sensitivity is the steady-state gain

$$S_I = \frac{p_3}{p_2}, \quad X_{ss} = S_I(I - I_b), \quad (5)$$

with units $\text{mL } \mu\text{U}^{-1}\text{min}^{-1}$. Parameters are mapped by sigmoid functions into physiologically sourced intervals informed by minimal-model studies in insulin-dependent diabetes [33] and regularized toward population anchors. Consequently, plausible range alone is not evidence of identifiability; external and stability checks are required.

Subcutaneous insulin absorption follows a two-compartment model [34]:

$$\dot{S}_1 = u_{\text{ins}}(t) - \frac{S_1}{t_{\max,I}}, \quad (6)$$

$$\dot{S}_2 = \frac{S_1}{t_{\max,I}} - \frac{S_2}{t_{\max,I}}. \quad (7)$$

For an impulse dose D , insulin on board is

$$S_1 + S_2 = D e^{-t/t_{\max,I}} \left(1 + \frac{t}{t_{\max,I}} \right), \quad (8)$$




Fig. 1. Audited digital-twin pipeline. OhioT1DM glucose and related patient data are aligned to five-minute intervals and processed through two parallel paths: a deep-learning model and a physiological model. The deep-learning model predicts future glucose levels and estimates patient-specific parameters, while the physiological model represents insulin and glucose dynamics. Physiological information is used both as an input to the model and as a training constraint to improve prediction consistency.

which starts at D and decreases monotonically. Carbohydrate absorption follows compartmental meal models used in T1D simulation [35], [36], [37] and is represented by stomach and gut states,

$$\dot{Q}_{\text{sto}} = -k_{\text{gri}}Q_{\text{sto}} + u_{\text{carb}}(t), \quad (9)$$

$$\dot{Q}_{\text{gut}} = k_{\text{gri}}Q_{\text{sto}} - k_{\text{abs}}Q_{\text{gut}}, \quad (10)$$

$$R_a(t) = f k_{\text{abs}}Q_{\text{gut}}(t), \quad (11)$$

so carbohydrate on board is $Q_{\text{sto}} + Q_{\text{gut}}$ and absorbed mass is conserved. For implementation, the non-glucose compartments are collected into

$$\mathbf{z} = [S_1, S_2, I, Q_{\text{sto}}, Q_{\text{gut}}, X]^\top, \quad \mathbf{u} = [u_{\text{ins}}, u_{\text{carb}}, I_b]^\top. \quad (12)$$

With parameters held constant over one grid interval, $\dot{\mathbf{z}} = A\mathbf{z} + B\mathbf{u}$ has the exact zero-order-hold update (a standard control-theory discretization technique, not attributed to a specific glucose-modeling source)

$$\mathbf{z}_{k+1} = A_d\mathbf{z}_k + B_d\mathbf{u}_k, \quad \begin{bmatrix} A_d & B_d \\ 0 & I \end{bmatrix} = \exp\left(\Delta \begin{bmatrix} A & B \\ 0 & 0 \end{bmatrix}\right). \quad (13)$$

This matrix-exponential formulation preserves the affine basal term through the explicit I_b input and avoids the accumulated integration error of a coarse explicit solver.

E. Feature Contract and Forecasting Network

Each timestep contains 35 named features in six groups: glucose level, multiscale rates of change and rolling statistics; delivered insulin and timing; meals and timing; mechanistic states (IOB, COB, plasma insulin, remote action, and glucose appearance); cyclic time; and context/sensor variables. A two-hour window $X \in \mathbb{R}^{24 \times 35}$ is standardized using training-only statistics.

The forecasting network is deliberately small (approximately 79,000 parameters) because the effective sample size is limited. A linear projection and positional encoding feed a Transformer encoder [38]. Attention pooling gives

$$z = \sum_{t=1}^{24} \alpha_t h_t, \quad \alpha_t = \frac{\exp(w^\top h_t)}{\sum_j \exp(w^\top h_j)}, \quad (14)$$

avoiding the equal weighting of old and recent measurements induced by mean pooling.

One head estimates the ten-component subject parameter vector θ_p . A second predicts a cubic B-spline trajectory

$$G_\theta(t) = G_0 + \sum_{k=1}^K c_k (B_k(t) - B_k(0)), \quad (15)$$

which anchors every trajectory at the last observed glucose G_0 . Median knot values are unconstrained, whereas lower and upper quantiles are generated by positive softplus offsets in glucose units; quantiles therefore cannot cross. A learned trust gate $g = \sigma(\gamma)$ controls the contribution of the mechanistic prior, with $g \approx 0.10$ at initialization. Predictions at all four horizons are evaluated from the spline.

F. Training Objective

The data objective combines Huber's robust loss [39] for the median with the pinball loss used in quantile regression [40] for lower and upper quantiles:

$$\mathcal{L}_{\text{data}} = \mathcal{L}_{\text{Huber}}(y, \hat{y}_{0.5}) + \sum_{q \in \{0.1, 0.9\}} \rho_q(y - \hat{y}_q), \quad (16)$$

where $\rho_q(u) = u(q - \mathbb{1}_{u < 0})$. Physics residuals are evaluated at 121 collocation points over the two-hour forecast and normalized by physiological scales, following the residual-minimization principle of PINNs [8]. The full objective is

$$\mathcal{L} = \mathcal{L}_{\text{data}} + \lambda_{\text{phys}} \mathcal{L}_{\text{phys}} + \lambda_{\text{prior}} \mathcal{L}_{\text{prior}} + \lambda_{\text{stab}} \mathcal{L}_{\text{within-subject}}. \quad (17)$$

Closed-form state updates replace a memory-intensive differentiable ODE solver. Adaptive loss weighting was motivated by work on gradient imbalance and uncertainty-based multi-objective learning [41], [42]. The selected A7 configuration applies physics from epoch one; a data-first curriculum was rejected because validation optima repeatedly occurred before the curriculum became selectable.

G. Evaluation and Statistical Analysis

The official temporal protocol uses each subject's earlier data for training. LOSO trains on 11 subjects and evaluates the held-out subject on the identical test windows.

Metrics are computed per subject and summarized as mean $\pm$ SD. Persistence predicts $\hat{G}(t+h) = G(t)$. Skill is $1 - \text{RMSE}_{\text{model}}/\text{RMSE}_{\text{persistence}}$.

Clinical evaluation includes Clarke and Parkes zones [43], [44], [45], consensus CGM ranges [52], coefficient-of-variation ratio, and Kovatchev low/high blood-glucose risk indices [46], [47], together with hypoglycemia sensitivity, specificity, precision, and quantile calibration. Paired model comparisons use the Wilcoxon signed-rank test [48] across 12 subjects with Holm multiplicity correction [49]; bootstrap intervals resample subjects, never overlapping windows. Test-retest reliability of S_I is measured with one-way ICC(1) [50], and concurrent validity uses Spearman rank correlation [51].

H. Preregistered Outcomes and Falsification Criteria

The analysis plan was fixed before the two experimental branches were evaluated. The primary outcome is 30-minute MAE under the official protocol for the physiology-focused branch. Secondary outcomes include all horizons under both protocols, persistence skill, Clarke and Parkes distributions, the full ablation matrix, quantile calibration, and the three S_I validation checks. Table II records the criteria that were fixed in advance; unmet criteria are reported as failures rather than redefined after inspection.

No asymmetric hypoglycemia/hyperglycemia point-loss term is permitted, because such a term would directly optimize the safety table later used as evidence. The test set is evaluated only after model selection, and the official protocol is never described as subject-disjoint generalization.

Integrated gradients (IG) [54] uses the standardized training mean as its zero baseline:

$$\text{IG}_i(x) = (x_i - x'_i) \int_0^1 \frac{\partial f(x' + \alpha(x - x'))}{\partial x_i} d\alpha. \quad (18)$$

Because IG shows an unresolved completeness violation, every interpretation is cross-checked with model-agnostic permutation importance, following the perturbation principle popularized for random forests [55].

III. RESULTS

A. Baseline Validation and Data Accounting

Persistence RMSE on the 2018 cohort is 22.60 ± 2.50 mg/dL at 30 minutes and 36.34 ± 3.14 mg/dL at 60 minutes, within 0.10 and 0.26 mg/dL, respectively, of two independently published no-change baselines [56], [57]. This close agreement validates the complete evaluation pipeline, including data parsing, five-minute grid alignment, horizon construction, persistence-target generation, RMSE computation, and subject-level aggregation Table III. The retained-window accounting further confirms that the reported metrics are computed on the intended 26,498 common test windows.

B. Official-Protocol Forecast Accuracy

Table IV and Fig. 2 show that the model improves on persistence at every horizon. All 12 subjects improve at 30 minutes; subject-level MAE skill ranges from 12.2% to 31.6%. The decline in R^2 with horizon is expected because future meals, boluses, and activity are absent from the input.

TABLE II
PREREGISTERED SUCCESS AND FALSIFICATION CRITERIA

Claim	Criterion fixed before test evaluation	Interpretation
Beats persistence	30-minute MAE below persistence in at least 10 of 12 subjects	Requires broadly distributed improvement, not a pooled gain driven by a few subjects.
Field-competitive point accuracy	30-minute MAE ≤ 13.5 mg/dL	A field-relative target, explicitly not a clinical acceptability threshold.
Physics improves accuracy	Physics arm beats A0 at 30 minutes and survives paired Wilcoxon testing with Holm correction	If unmet, physics is not credited with point-accuracy improvement.
S_I is subject-specific	Between-subject CV $> 10\%$	Rejects a parameter head collapsed to one population value.
S_I is stable	ICC(1) > 0.5 on disjoint within-subject blocks	Rejects a window-level nuisance variable that merely absorbs prediction error.
S_I is externally consistent	Spearman $\rho < 0$ versus total daily dose, with a subject-bootstrap 95% CI excluding zero	Evaluated separately under personalized and LOSO protocols.




Fig. 2. Official-protocol accuracy summary at 30/60/90/120 minutes. Left and center: MAE and RMSE for persistence versus the digital twin both rise with horizon, but the gap between the two curves widens rather than closes, showing the model’s advantage over naive persistence is not a short-horizon artifact. Right: the resulting RMSE skill stays essentially flat at 19–21% across all four horizons. Error bars are subject-level SD across the 12 OhioT1DM subjects.

TABLE III
WINDOW ACCOUNTING AFTER INTEGRITY FILTERING

Quantity	Value
Subjects / CGM observations	12 / 166,443
Candidate windows	187,780
Retained windows	141,100 (75.1%)
Incomplete input span	36,395
Missing real target	10,285
Common test windows	26,498
Effective independent windows	$\approx 2,800$

C. Cross-Subject Generalization

LOSO remains better than persistence in 10 of 12 subjects at 30 minutes, but is consistently worse than the personalized protocol. As Fig. 3 shows, the personalization gap grows with horizon, consistent with increasing dependence on individual insulin kinetics and routine.

D. Physics Ablations

No physics configuration significantly outperforms the no-physics A0 baseline after Holm correction (Table V; Fig. 4). A7 is nominally best and improves 9 of 12 subjects, but its corrected $p = 0.105$. Per-patient parameters are accuracy-neutral: A3 and population-fixed A4 differ by 0.07–0.23 mg/dL at 30/60 min, widening to 0.76 mg/dL at 120 min. In contrast with earlier physiology-informed OhioT1DM studies that primarily reported forecast accuracy [58], [59], this ablation separates parameter estimation from predictive gain. The




Fig. 3. Matched-protocol comparison across 30, 60, 90, and 120-minute horizons. Persistence, personalized, and LOSO models are evaluated on identical windows. LOSO closely follows the personalized model at 30 minutes but increasingly approaches persistence at longer horizons, indicating a greater benefit from within-subject history for extended forecasting.

defensible value of physics is not superior point accuracy; it is access to a stable physiological parameter and informative mechanistic states without a material accuracy penalty.

E. Clinical Accuracy and Hypoglycemia Detection

Clinical agreement remains high but decreases with forecast horizon: Clarke A+B decreases from 99.0% at 30 minutes to 94.3% at 120 minutes (Fig. 5). However, Clarke zone A mea-

TABLE IV
OFFICIAL PERSONALIZED PROTOCOL: MEAN $\pm$ SD ACROSS SUBJECTS (MG/DL UNLESS NOTED)

Horizon	Persistence MAE	Model MAE	Persistence RMSE	Model RMSE	RMSE skill	R^2	MARD
30 min	16.87 $\pm$ 1.92	13.08 $\pm$ 1.73	23.37	18.84 $\pm$ 2.58	19.4%	0.885	8.79%
60 min	28.22 $\pm$ 3.24	21.98 $\pm$ 2.96	38.15	30.52 $\pm$ 3.98	20.0%	0.700	14.76%
90 min	36.55 $\pm$ 4.00	28.42 $\pm$ 3.89	48.69	38.60 $\pm$ 5.05	20.7%	0.523	19.14%
120 min	42.90 $\pm$ 4.46	33.23 $\pm$ 4.45	56.43	44.46 $\pm$ 5.62	21.2%	0.368	22.37%

TABLE V
ABLATION MATRIX. HYPOGLYCEMIA BIAS IS MEAN ERROR FOR ACTUAL GLUCOSE BELOW 70 MG/DL

Arm	Configuration	MAE30	MAE60	MAE120	Clarke A	Hypo sens.	Hypo bias	Gives S_I
A0	Transformer, no physics	13.51	22.82	34.93	89.4%	0.637	+7.29	no
A1	Penalty PINN, fixed $\lambda = 0.1$	13.68	23.47	36.22	88.9%	0.502	+11.29	no
A2	Adaptive penalty, per patient	13.58	23.35	35.74	89.0%	0.499	+11.03	yes
A3	Hybrid prior, curriculum	13.65	23.12	34.83	89.0%	0.485	+10.95	yes
A4	Hybrid, fixed population parameters	13.58	23.35	35.59	89.1%	0.517	+10.18	no
A7	Hybrid, full physics from epoch 1	13.25	22.29	33.38	89.8%	0.608	+8.61	yes




Fig. 4. Point-accuracy ablation at 30 minutes comparing eight physics configurations (A0-A7), from no physics (A0) to full physics from epoch 1 (A7). A7 has the lowest MAE (13.25 mg/dL) versus 13.51 mg/dL for A0; however, the small 0.26 mg/dL difference is not statistically significant after Holm correction (Table V), indicating no clear accuracy advantage of any physics configuration.

sure point-wise clinical agreement rather than hypoglycemia-event detection. Thus, 89.8% of predictions falling in zone A and 0% in zone E at 30 minutes can still coexist with the median forecast missing 43% of hypoglycemic events.

The lower predictive quantile is well calibrated, with observed coverage of 0.095 for nominal $q = 0.10$ and 0.889 for nominal $q = 0.90$. Using $\hat{G}_{0.10} < 70$ mg/dL as a hypoglycemia alarm substantially increases sensitivity from 0.566 to 0.928 (Fig. 6), improving detection for all eight subjects with at least 20 hypoglycemic events. This improvement comes at a clear cost: precision decreases from 0.699 to 0.347, while false positives increase from 173 to 1,257. Thus, the quantile-based alarm represents a tunable sensitivity-precision operating point rather than evidence of deployment readiness.

Point forecasts also increasingly compress glucose excursions at longer horizons. The predicted-to-actual coefficient-of-variation ratio decreases from 0.973 to 0.802, with TIR increasingly overestimated and TBR underestimated (Fig. 7). Consistently, predicted LBGI falls increasingly below the




Fig. 5. Clarke error-grid composition across 30/60/90/120-minute horizons. Zone A decreases from $\sim 90\%$ to $\sim 58\%$, with the difference shifting mainly to zone B, while dangerous zones D and E remain negligible. Despite this favorable composition, the median forecast still misses many hypoglycemic events, so low D/E rates should not be interpreted as high hypoglycemia sensitivity.

actual value, from 0.65 versus 0.77 at 30 minutes to 0.33 versus 0.74 at 120 minutes, indicating progressively weaker representation of hypoglycemic risk in longer-horizon point forecasts.

F. Patient-Specific Insulin Sensitivity

Three preregistered checks determine whether S_I may be interpreted physiologically (Table VI). All pass under the official protocol: between-subject variation is non-degenerate, test-retest reliability is high, and S_I is strongly inversely related to total daily insulin dose. Under LOSO, stability and variability remain, but the dose correlation fails. Hence S_I is a stable, subject-specific estimate with external validity demonstrated only in the personalized setting; cross-subject validity is not established.

Body weight is not identifiable because all OhioT1DM files use the same de-identification placeholder. The dose-per-




Fig. 6. Hypoglycemia alarm trade-off at 30 minutes comparing the median forecast with the $q = 0.10$ lower-quantile alarm. The quantile alarm increases sensitivity from 0.566 to 0.928 with little change in specificity, but reduces precision from 0.699 to 0.347 due to more frequent false alarms. This represents a tunable sensitivity-precision trade-off, not deployment readiness.

TABLE VI
PREREGISTERED VALIDATION OF INSULIN SENSITIVITY

Check	Threshold	Official	LOSO
Between-subject CV	$>10\%$	35.8%	34.3%
Test-retest ICC(1)	>0.5	0.890	0.816
ρ vs. total daily dose	<0 , CI excludes 0	-0.839	-0.357
95% CI for ρ	-	[-0.99, -0.44]	[-0.82, 0.28]

kilogram correlation is therefore not independent of total daily dose. Moreover, p_2 and p_3 are range-constrained, so agreement with a plausible interval is partly imposed rather than learned.

G. Feature Attribution

IG and permutation importance agree with Spearman $\rho = 0.731$ ($p = 6.05 \times 10^{-7}$, $n = 35$) and share three of their top five features. Both rank five-minute glucose rate of change first by a wide margin. Permuting it raises MAE by 3.11 mg/dL, over seven times the next feature. The group-level IG result in Fig. 8 assigns 42.7% to glucose-derived features and 27.1% to mechanistic states, showing that physiology-derived inputs are informative even though the physics loss does not significantly improve point accuracy.

Attribution concentrates in the most recent input positions, supporting attention pooling. The IG completeness discrepancy has median relative magnitude 3.2% and p_{95} 70%; it does not shrink with more integration steps and remains unresolved. Conclusions are therefore restricted to findings corroborated by permutation importance.

H. Comparison With Published Work

At 30 minutes, the current work is mid-field relative to the 2020 challenge literature [60], [61], [62], [63], [64] (Table VII): RMSE is 18.84 and MAE is 13.08 mg/dL. At 60 minutes, RMSE 30.52 and MAE 21.98 are nominally better than the verified comparison set, but differences in window eligibility, missing-data handling, external pretraining, and best-configuration reporting prevent a state-of-the-art claim. The

defensible advances are calibrated hypoglycemia detection, explicit personalized-versus-LOSO decomposition, preregistered validation of a physiological parameter, and a quantified account of the benefits and costs of Bergman constraints. The Current-Work figures quoted here (18.84/13.08/30.52/21.98) are the official, all-12-subject personalized-protocol numbers reported throughout the Results section; Table I instead reports the slightly different 2020-cohort-matched subset (18.74/13.24/31.16/22.55, subjects 540/544/552/567/584/596 only), because that is the exact 6-subject cohort the 2020 challenge entries were themselves scored on. Both numbers are correct; they answer different comparability questions and should not be averaged together.

IV. DISCUSSION

The results separate three concepts that are often conflated in digital-twin and glucose-prediction literature [2]. First, a mechanistic constraint need not improve point accuracy to be useful. Here it supplies interpretable state variables and a stable S_I estimate at little accuracy cost, while mechanistic features account for more than one quarter of total IG attribution. Second, a favorable error grid does not demonstrate event safety. The median forecast has excellent Clarke A+B coverage but inadequate hypoglycemia sensitivity; a calibrated predictive tail is required. Third, personalized temporal validation and cross-subject generalization answer different questions. The growing official-LOSO gap shows that access to a subject's history matters increasingly at longer horizons. Prospective digital-twin trials demonstrate the longer-term potential of individualized replay and adaptation, but do not remove the need to validate the forecasting layer independently [18].

The ablation also exposes a structural limitation. Bergman's equilibrium at G_b pulls trajectories toward the center of the glucose distribution. In the best configuration this adds approximately 1.3 mg/dL to an already-present regression-to-the-mean bias during hypoglycemia. This behavior appears whether physics enters as a penalty or a prior. A future model should modify the physiological structure rather than hide the bias with an asymmetric point loss.

V. LIMITATIONS AND FUTURE WORK

The cohort contains only 12 subjects and approximately 2,800 effectively independent observations. All reported runs use one random seed. OhioT1DM has no clamp measurement of S_I , so validation is indirect; LOSO external validity fails. Subject 552 has only 59.7% test CGM coverage, and subject 567 has no carbohydrate records in the test interval. Body weight is unavailable. Quantile evaluation reuses the ablation test set, and the alarm operating point lacks an elicited clinical cost ratio. No counterfactual intervention is validated, so the system is a forecasting-centered digital twin rather than a validated treatment simulator.

The parameter head estimates the full ten-component physiological vector θ_p per subject (Table IX: $p_1, p_2, p_3, n, V_G, V_I, t_{\max, I}, k_{\text{gri}}, k_{\text{abs}}, f$), not only S_I . Preregistered external validation (Table VI) was carried




Fig. 7. Excursion compression with horizon (30/60/90/120 minutes). Left: actual time-in-range (TIR) stays flat near 62% while predicted TIR climbs toward 74%, an increasing overestimate. Center: actual time-below-range (TBR) stays flat near 2.8% while predicted TBR falls toward 0.6%, an increasing underestimate of hypoglycemic exposure. Right: the predicted-to-actual coefficient-of-variation ratio declines from about 0.97 to 0.80. All three panels show the same mechanism: as the horizon lengthens, the point forecast increasingly regresses toward the mean and loses the variance of the true glucose trace.




Fig. 8. Feature attribution at 30 minutes from two independent methods. Left: integrated-gradient (IG) share of total attribution by feature group – glucose-derived features dominate (42.7%), followed by the mechanistic states (27.1%) derived from the Bergman-plus-absorption model, with therapy, time, sensor, and context features contributing the remainder. Right: model-agnostic permutation importance, ranked by MAE increase when a feature is shuffled; the five-minute glucose rate of change (roc_5) dominates at roughly 3.1 mg/dL, more than seven times the next-largest feature, with the remaining features forming a long, comparatively flat tail. The two methods agree on the leading features despite being computed independently.

TABLE VII
PUBLISHED OHIO T1DM COMPARISON (MG/DL)

Method	R30	A30	R60	A60
Freiburghaus CNN/LSTM [63]	17.45	11.22	33.67	23.25
Rubin-Falcone N-BEATS+BiLSTM [62]	18.22	12.83	31.66	23.60
Bevan-Coenen LSTM [61]	18.23	14.37	31.10	25.75
Pavan NN-EIM [60]	18.63	10.08	32.27	17.69
Yang MS-LSTM [64]	19.05	13.50	32.03	23.83
Current Work, official	18.84	13.08	30.52	21.98
Current Work, LOSO	20.73	14.66	33.96	24.85
Persistence	22.50	16.87	36.60	28.22

out for S_I alone, for three concrete reasons rather than by omission. First, S_I is the one component with a direct clinical measurement analogue: a euglycemic clamp measures exactly this quantity, giving a real external target to validate against, whereas the remaining parameters have no equivalent bedside measurement available in OhioT1DM. Second, S_I has literature-sourced plausible bounds specific to insulin-dependent diabetes [33], giving its sigmoid-mapped range a population anchor that is not equally well established for every other compartment parameter. Third, the ablation (Table V, A3 vs. A4) already shows that per-patient parameter estimation is accuracy-neutral overall (0.07-0.76 mg/dL), so

the contribution claimed here is the external validity of one physiological quantity, not evidence that the other nine are unimportant. One concrete obstacle already blocks part of this extension: body weight is unavailable because every OhioT1DM file carries the same de-identification placeholder, which is why a body-weight-normalized analysis was not attempted for any parameter, not only S_I .

Priority extensions are multi-seed replication, external validation on a larger cohort, identifiability analysis of patient parameters, preregistered validation of the remaining Bergman and absorption parameters (p_1, n, V_G, V_I) against literature-sourced population anchors comparable to those used for

S_I , time-varying S_I , a joint multi-horizon predictive distribution, decision-theoretic alarm selection, longer context, and prospective validation of counterfactual meal and insulin interventions. Prediction Error-Grid Analysis (PRED-EGA) was designed specifically for glucose predictors [53] and is a natural future addition, but it is not reported here because its complete operational boundaries were not independently reconstructed in the audited implementation.

VI. CONCLUSION

An audited, physics-guided digital twin can outperform validated persistence across 30–120-minute horizons while making its limitations visible. The strongest evidence is not that physics improves point accuracy, it does not do so significantly but that physiology can provide stable patient-specific structure at negligible predictive cost and that calibrated distributional forecasting can recover hypoglycemic events missed by an otherwise accurate median forecast. Horizon integrity, subject-level statistics, matched generalization protocols, and explicit safety trade-offs are prerequisites for credible glucose forecasting.

APPENDIX A NOTATION, VARIABLES, AND TERMINOLOGY

This appendix defines the symbols, terms, abbreviations, and feature groups used in the manuscript. Rate constants are expressed per minute unless stated otherwise. A subscript b denotes a basal or fasting quantity, a hat denotes a prediction, bold lowercase symbols denote vectors, and bold uppercase symbols denote matrices.

A. Physiological States and Inputs

Glucose G is propagated separately from a six-state linear compartment model. Insulin on board and carbohydrate on board are derived identities, not additional fitted states (cf. Eqs. (6)–(7), (9)–(10)):

$$\text{IOB} = S_1 + S_2, \quad \text{COB} = Q_{\text{sto}} + Q_{\text{gut}}. \quad (19)$$

B. Estimated Physiological Parameters

The parameter head emits unconstrained values z_i , which are mapped to admissible, physiologically sourced intervals [33] by

$$\theta_i = \ell_i + (u_i - \ell_i)\sigma(z_i). \quad (20)$$

The bounds prevent nonphysical values but also mean that plausible range alone cannot establish identifiability.

The reported derived quantity is

$$S_I = p_3/p_2, \quad (21)$$

the steady-state gain from plasma insulin above basal to remote insulin action, already introduced as Eq. (5) [1], [32]. Higher S_I means a given insulin excess produces stronger glucose disposal. It is validated as a subject-specific estimate only when its cross-subject variation, test–retest ICC, and dose correlation meet the preregistered criteria.

The implementation-level objective, following Eq. (17), can be written more explicitly as

$$\mathcal{L} = \mathcal{L}_{\text{Huber}} + \lambda_q \mathcal{L}_{\text{pinball}} + w_{\text{phys}} \mathcal{L}_{\text{res}} + \lambda_{\text{pr}} \mathcal{L}_{\text{prior}} + \lambda_{\text{tc}} \mathcal{L}_{\text{temporal}}. \quad (22)$$

C. Feature Contract

The 35-input feature contract contains six groups: nine glucose features (current glucose, three rates of change, and rolling summaries); nine mechanistic features (IOB, COB, the six compartment states, and glucose appearance); four therapy/timing features; three cyclic-time features; six sensor values or availability masks; and four behavioral/context indicators. Glucose features account for 42.7% of IG magnitude, mechanistic features 27.1%, therapy 13.5%, time 8.1%, sensors 7.0%, and context 1.7%. Every sensor channel has an availability indicator so missingness is not represented as a physiologically extreme standardized zero. Input interpolation is explicitly flagged, whereas forecast targets are always real CGM observations.

For reference, Table VIII summarizes physiological state and input notation, Table IX the estimated patient-parameter vector, Table X architecture/discretization/uncertainty notation, Table XI training-loss terminology, Table XII evaluation-metric definitions, and Table XIII the abbreviations used throughout this manuscript.

TABLE VIII
PHYSIOLOGICAL STATE VARIABLES AND EXTERNAL INPUTS

Symbol	Unit	Name	Meaning in this study
$G(t)$	mg/dL	Plasma glucose	Glucose concentration forecast by the learned spline and constrained by the mechanistic equation.
G_b	mg/dL	Basal glucose	Subject-specific fasting anchor used by the Bergman equilibrium.
G_0	mg/dL	Anchor glucose	Last real CGM observation at forecast time $t = 0$.
$X(t)$	min^{-1}	Remote insulin action	Interstitial insulin effect entering glucose disposal additively with p_1 .
$I(t)$	$\mu\text{U/mL}$	Plasma insulin	Circulating insulin concentration generated by subcutaneous absorption and clearance.
I_b	$\mu\text{U/mL}$	Basal plasma insulin	Concentration implied by the subject's basal delivery; subtracting it preserves basal equilibrium.
S_1, S_2	U	SC insulin states	First and second subcutaneous absorption compartments.
Q_{sto}	g	Stomach carbohydrate	Ingested carbohydrate awaiting gastric emptying.
Q_{gut}	g	Gut carbohydrate	Carbohydrate available for intestinal absorption.
$R_a(t)$	mg/min	Glucose appearance rate	Meal-derived glucose entering the accessible glucose compartment.
u_{ins}	U/min	Insulin input	Basal plus bolus insulin delivery on the five-minute grid.
u_{carb}	g/min	Meal input	Recorded carbohydrate ingestion on the grid.
$\mathbf{z}$	–	Compartment state	$[S_1, S_2, I, Q_{\text{sto}}, Q_{\text{gut}}, X]^T$.
$\mathbf{u}$	–	Input vector	$[u_{\text{ins}}, u_{\text{carb}}, I_b]^T$.

TABLE IX
PATIENT-PARAMETER VECTOR θ_p AND ADMISSIBLE INTERVALS

Symbol	Meaning	Range	Unit	Population anchor
p_1	Insulin-independent glucose disposal	0–0.030	min^{-1}	0.013
p_2	Decay of remote insulin action	0.005–0.10	min^{-1}	0.025
p_3	Build rate of action from insulin excess	10^{-6} – 3×10^{-5}	$\text{mL } \mu\text{U}^{-1} \text{min}^{-2}$	6.25×10^{-6}
n	Plasma insulin clearance	0.08–0.25	min^{-1}	0.138
V_G	Glucose distribution volume	1.4–2.4	dL/kg	1.88
V_I	Insulin distribution volume	0.08–0.18	L/kg	0.12
$t_{\text{max}, I}$	Time to peak SC absorption	30–90	min	55
k_{gri}	Gastric emptying rate	0.008–0.10	min^{-1}	0.035
k_{abs}	Intestinal absorption rate	0.005–0.10	min^{-1}	0.0167
f	Carbohydrate bioavailability	0.70–1.00	dimensionless	0.90

TABLE X
ARCHITECTURE, DISCRETIZATION, AND UNCERTAINTY NOTATION

Symbol	Value or dimension	Definition
t	0–120 min	Time measured from the forecast anchor.
h	30, 60, 90, 120 min	Forecast horizon.
Δ	5 min	Common temporal grid step.
B	batch dependent	Number of windows processed together.
H	4	Number of reported forecast horizons.
K	12	Number of cubic B-spline basis functions.
$B_k(t)$	spline basis	The k th clamped, uniformly knotted B-spline basis function.
c_k	learned scalar	Spline coefficient emitted by the trajectory head.
$G_\theta(t)$	mg/dL	Anchored learned trajectory $G_0 + \sum_k c_k [B_k(t) - B_k(0)]$.
A, B	$6 \times 6, 6 \times 3$	Continuous compartment-state and input matrices.
A_d, B_d	discrete matrices	Exact zero-order-hold matrices obtained from one augmented matrix exponential.
θ_p	10 elements	Patient-specific physiological parameter vector in Table IX.
$g = \sigma(\gamma)$	initialized near 0.10	Learned trust gate controlling the mechanistic prior.
$\hat{G}_{0.5}$	mg/dL	Median forecast and reported point prediction.
$\hat{G}_q$	mg/dL	Forecast at quantile level $q \in \{0.1, 0.5, 0.9\}$.
$\delta_q(h)$	mg/dL	Positive softplus-derived distance from the median, preventing quantile crossing.
$r(t)$	dimensionless	Scale-normalized physics residual evaluated at 121 collocation points.
n_{eff}	$\approx N/48$	Approximate independent sample count after accounting for window overlap.

TABLE XI
TRAINING-LOSS TERMINOLOGY

Term	Meaning
$\mathcal{L}_{\text{Huber}}$	Horizon-weighted robust point loss for the median forecast; quadratic near zero error and linear for large errors.
$\mathcal{L}_{\text{pinball}}$	Proper quantile loss $\max\{q(y - \hat{y}_q), (q - 1)(y - \hat{y}_q)\}$.
$\mathcal{L}_{\text{res}}$	Mean squared, nondimensionalized Bergman residual over the collocation grid.
$\mathcal{L}_{\text{prior}}$	Weak population-prior penalty that discourages physiologically implausible parameter drift.
$\mathcal{L}_{\text{temporal}}$	Within-subject penalty on disagreement between parameter estimates from adjacent windows.
w_{phys}	Physics weight, adaptive in the relevant ablation arms and optionally multiplied by a curriculum ramp.
$s_i = \log \sigma_i^2$	Stable log-variance parameterization used for learned multi-objective weighting [42].

TABLE XII
EVALUATION METRICS AND CLINICAL TERMS

Term	Unit	Definition and interpretation
MAE	mg/dL	Mean absolute error, $n^{-1} \sum_i \hat{y}_i - y_i $.
RMSE	mg/dL	Root-mean-square error, emphasizing larger deviations.
R^2	—	$1 - \text{SS}_{\text{res}}/\text{SS}_{\text{tot}}$.
MARD	%	Mean absolute relative difference, with reference glucose in the denominator.
Skill	%	$1 - \text{RMSE}_{\text{model}}/\text{RMSE}_{\text{persistence}}$. Positive values beat no-change prediction.
TIR/TAR/TBR	%	Time in, above, or below the consensus glucose range. The main in-range interval is 70–180 mg/dL [52].
CV ratio	—	$\text{CV}_{\text{pred}}/\text{CV}_{\text{actual}}$; values below one quantify excursion compression.
LBGI/HBGI	—	Low/high blood-glucose risk indices after the Kovatchev symmetric risk transform.
Coverage	—	Fraction of observations below a stated predictive quantile; calibration requires observed and nominal coverage to agree.
Hypo sensitivity	—	Fraction of actual $G < 70$ mg/dL events detected by the forecast or alarm.
Precision	—	Fraction of hypoglycemia alarms corresponding to actual events.
ICC(1)	—	One-way, single-measure intraclass correlation used for S_I test–retest stability [50].
ρ	—	Spearman rank correlation used for monotonic association with insulin dose [51].
IG_i	mg/dL	Integrated-gradient attribution of feature i relative to the standardized-mean baseline.
Clarke/Parkes EGA	zone	Point-wise clinical error grids originally designed for glucose-measurement error.
PRED-EGA	category	Prediction-specific error grid combining level and rate information [53]; not implemented here.

TABLE XIII
ABBREVIATIONS AND PROTOCOL TERMINOLOGY

Short form	Expansion	Study-specific meaning
T1D/IDDM	Type 1 diabetes / insulin-dependent diabetes mellitus	IDDM is retained only when describing older physiological literature.
CGM	Continuous glucose monitoring	Nominal five-minute glucose sampling in OhioT1DM.
SC	Subcutaneous	Route represented by insulin compartments S_1 and S_2 .
IOB/COB	Insulin/carbohydrate on board	Derived as $S_1 + S_2$ and $Q_{\text{sto}} + Q_{\text{gut}}$.
TDD/ICR	Total daily insulin dose / insulin-to-carbohydrate ratio	External consistency variables for the physiology analysis; ICR is available only for six subjects.
PINN/ODE	Physics-informed neural network / ordinary differential equation	Neural training constrained by the physiological differential equations.
LOSO	Leave-one-subject-out	Twelve subject-disjoint folds, each withholding all data from one test subject.
Official protocol	OhioT1DM temporal holdout	Later data from subjects whose earlier data are available for training; therefore personalized.
BGLP	Blood Glucose Level Prediction challenge	OhioT1DM challenge setting; 2018 and 2020 have different test warmup rules.
EGA	Error-grid analysis	Clarke and Parkes clinical zone classification.
IG/SHAP	Integrated gradients / SHapley Additive exPlanations	IG is used; the legacy single-timestep SHAP procedure was rejected.
SD/CI/ICC	Standard deviation / confidence interval / intraclass correlation	Uncertainty and reliability summaries computed at the subject level.
A0–A7	Ablation arms	Controlled variants of physics form, weighting, parameter personalization, and curriculum.
ZOH	Zero-order hold	Exact interval propagation under piecewise-constant inputs; also used in the literature for persistence.

REFERENCES

- [1] R. N. Bergman, Y. Z. Ider, C. R. Bowden, and C. Cobelli, "Quantitative estimation of insulin sensitivity," *Am. J. Physiol.*, vol. 236, no. 6, pp. E667–E677, 1979.
- [2] G. Cappon and A. Facchinetti, "Digital twins in type 1 diabetes: A systematic review," *J. Diabetes Sci. Technol.*, vol. 19, no. 1, 2025.
- [3] C. Duckworth, M. J. Guy, *et al.*, "Explainable machine learning for real-time hypoglycemia and hyperglycemia prediction and personalized control recommendations," *J. Diabetes Sci. Technol.*, vol. 18, no. 1, 2024.
- [4] R. Hasan, V. Dattana, S. Mahmood, and S. Hassan, "Towards transparent diabetes prediction: combining AutoML and explainable AI for improved clinical insights," *Information*, vol. 16, no. 1, p. 7, 2025.
- [5] S. L. Cichosz, S. S. Olesen, and M. H. Jensen, "Explainable machine-learning models to predict weekly risk of hyperglycemia, hypoglycemia, and glycemic variability in patients with type 1 diabetes based on continuous glucose monitoring," *J. Diabetes Sci. Technol.*, vol. 20, no. 3, pp. 836–847, 2024.
- [6] A. Adjevi, A. M. Abdirashid, F. Aktaş, M. H. B. Ucar, and S. Solak, "Explainable reinforcement learning for glucose monitoring based on Shapley value analysis," *Comput. Methods Programs Biomed.*, vol. 278, 2026.
- [7] A. Y. Ahmed, "Personalized glycemic control modeling using explainable artificial intelligence for optimizing diabetes management," *Scientia. Technology, Science and Society*, vol. 3, no. 1, 2026.
- [8] M. Raissi, P. Perdikaris, and G. E. Karniadakis, "Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations," *J. Comput. Phys.*, vol. 378, pp. 686–707, 2019.
- [9] R. Pei and W. Wang, "PINN based short-term glucose prediction for type 1 diabetes," in *2025 4th Asia Conf. Algorithms, Computing and Machine Learning (CACML)*, Guangzhou, China, 2025, pp. 1–6.
- [10] M. Wang, H. Zhang, and R. Song, "Blood glucose prediction algorithm via PINN under Bergman's Minimal Model constraints," in *2025 IEEE 14th Data Driven Control and Learning Systems Conf. (DDCLS)*, Wuxi, China, 2025, pp. 1797–1802.
- [11] W. Wang, R. Pei, D. Li, S. Liu, Y. Geng, and S. Wang, "A physics-informed glucose-insulin neural network model for glucose prediction," *Tsinghua Sci. Technol.*, 2026.
- [12] P. A. Mongini, M. Drecogna, and C. Toffanin, "New physics-LSTM hybrid models for control-oriented glucose prediction in Type 1 Diabetes," in *2025 American Control Conf. (ACC)*, Denver, CO, USA, 2025, pp. 662–667. [Verified against the corresponding author's institutional CV.]
- [13] P. Colmegna, K. Wang, J. Garcia-Tirado, and M. D. Breton, "Mapping data to virtual patients in type 1 diabetes," *Control Eng. Pract.*, vol. 103, Art. no. 104605, 2020, doi: 10.1016/j.conengprac.2020.104605. [Record verified via Crossref; content-match against the citing sentence not fully certain — see verification report.]
- [14] G. Cappon, M. Vettoretti, G. Sparacino, S. Del Favero, and A. Facchinetti, "ReplayBG: A digital twin-based methodology to identify a personalized model from type 1 diabetes data and simulate glucose concentrations to assess alternative therapies," *IEEE Trans. Biomed. Eng.*, vol. 70, no. 11, pp. 3227–3238, 2023, doi: 10.1109/TBME.2023.3286856. [Verified via Crossref.]
- [15] M. Hwang, V. P. Rachim, J. Yoo, Y. Lee, and S.-M. Park, "Generalized multi task learning framework for glucose forecasting and hypoglycemia detection using simulation to reality," *npj Digit. Med.*, vol. 8, Art. no. 612, 2025, doi: 10.1038/s41746-025-01994-4.
- [16] V. Roquemen-Echeverri, T. Kushner, P. G. Jacobs, and C. Mosquera-Lopez, "A physiologically-constrained neural network digital twin framework for replicating glucose dynamics in type 1 diabetes," *Neural Comput. Appl.*, vol. 38, p. 241, 2026.
- [17] C. E. Builes-Montañó, L. Lema-Perez, A. Ramírez-Rincón, J. J. Zuleta-Tobón, J. C. Restrepo-Gutiérrez, H. D. Álvarez-Zapata, J. García-Tirado, *et al.*, "A digital twin-enhanced decision support system improves time-in-range in type 1 diabetes: a randomised clinical trial," *Sci. Rep.*, vol. 15, Art. no. 39738, 2025, doi: 10.1038/s41598-025-23165-x. [Verified: content match confirmed.]
- [18] B. Kovatchev, P. Colmegna, J. Pavan, J. Diaz Castañeda, M. Villa Tamayo, C. Koravi, G. Santini, C. Alix, M. Stumpf, and S. Brown, "Human-machine co-adaptation to automated insulin delivery: A randomised clinical trial using digital twin technology," *npj Digit. Med.*, vol. 8, 2025.
- [19] E. Acuña and R. Aparicio, "Predicting the blood glucose level using transformers," in *Proc. 2023 Int. Conf. Computational Science and Computational Intelligence (CSCI)*, 2023, pp. 1392–1399, doi: 10.1109/CSCI62032.2023.00228.
- [20] A. Marx, F. Di Stefano, H. Leutheuser, K. Chin-Cheong, M. Pfister, M.-A. Burckhardt, S. Bachmann, and J. E. Vogt, "Blood glucose forecasting from temporal and static information in children with T1D," *Front. Pediatr.*, vol. 11, Art. no. 1296904, 2023, doi: 10.3389/fped.2023.1296904.
- [21] Q. Bian, A. As'arry, X. Cong, K. A. Md Rezali, and R. M. K. Raja Ahmad, "A hybrid Transformer-LSTM model applied to glucose prediction," *PLOS ONE*, vol. 19, no. 9, Art. no. e0310084, 2024.
- [22] X. Xiong, X. Yang, Y. Cai, Y. Xue, J. He, and H. Su, "Exploring the potential of deep learning models integrating transformer and LSTM in predicting blood glucose levels for T1D patients," *Digit. Health*, vol. 11, Art. no. 20552076251328980, 2025, doi: 10.1177/20552076251328980.
- [23] F. J. Lara-Abelenda, D. Chushig-Muzo, P. Peiro-Corbacho, A. M. Wägner, C. Granja, and C. Soguero-Ruiz, "Personalized glucose forecasting for people with type 1 diabetes using large language models," *Comput. Methods Programs Biomed.*, vol. 265, Art. no. 108737, 2025, doi: 10.1016/j.cmpb.2025.108737.
- [24] D. Darpit, K. Vyas, J. K. Jayagopal, A. Garcia, M. Erraguntla, and M. Lawley, "A personalized federated learning-based glucose prediction algorithm for high-risk glycemic excursion regions in type 1 diabetes," *Sci. Rep.*, vol. 15, Art. no. 38376, 2025.
- [25] M. P. Barbato, G. Rigamonti, D. Marelli, and P. Napoletano, "Lightweight sequential transformers for blood glucose level prediction in type-1 diabetes," *IEEE J. Biomed. Health Inform.*, vol. 30, no. 6, pp. 5465–5475, 2026, doi: 10.1109/JBHI.2025.3633194.
- [26] A. J. Rodriguez-Almeida, C. Betancort, A. M. Wägner, G. M. Callico, and H. Fabelo, "Incorporating uncertainty estimation and interpretability in personalized glucose prediction using the Temporal Fusion Transformer," *Sensors*, vol. 25, no. 15, Art. no. 4647, 2025.

- [27] Y. Lu *et al.*, “A pretrained transformer model for decoding individual glucose dynamics from continuous glucose monitoring data,” *Natl. Sci. Rev.*, vol. 12, no. 5, Art. no. nwaf039, 2025.
- [28] D. Wang, J. Liang, J. Ye, J. Li, J. Li, Q. Zhang, Q. Hu, C. Pan, D. Wang, Z. Liu, W. Shi, D. Shi, F. Li, B. Qu, and Y. Zheng, “Enhancement of the performance of large language models in diabetes education through retrieval-augmented generation: comparative study,” *J. Med. Internet Res.*, vol. 26, Art. no. e58041, 2024. [Author list verified.]
- [29] S. B. Soumna and H. Ghasemzadeh, “GlyRAG: Context-aware retrieval-augmented framework for blood glucose forecasting,” arXiv preprint, 2026.
- [30] C. Marling and R. Bunescu, “The OhioT1DM dataset for blood glucose level prediction,” in *Proc. 3rd Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2148, pp. 60–63, 2018.
- [31] C. Marling and R. Bunescu, “The OhioT1DM dataset for blood glucose level prediction: Update 2020,” in *Proc. 5th Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2675, pp. 71–74, 2020.
- [32] R. N. Bergman, L. S. Phillips, and C. Cobelli, “Physiologic evaluation of factors controlling glucose tolerance in man,” *J. Clin. Invest.*, vol. 68, no. 6, pp. 1456–1467, 1981.
- [33] G. M. Ward *et al.*, “A modified minimal model analysis of insulin sensitivity and glucose-mediated glucose disposal in insulin-dependent diabetes,” *Metabolism*, vol. 40, no. 1, pp. 4–9, 1991.
- [34] R. Hovorka *et al.*, “Nonlinear model predictive control of glucose concentration in subjects with type 1 diabetes,” *Physiol. Meas.*, vol. 25, no. 4, pp. 905–920, 2004.
- [35] E. D. Lehmann and T. Deutsch, “A physiological model of glucose–insulin interaction in type 1 diabetes mellitus,” *J. Biomed. Eng.*, vol. 14, no. 3, pp. 235–242, 1992.
- [36] C. Dalla Man, M. Camilleri, and C. Cobelli, “A system model of oral glucose absorption: Validation on gold standard data,” *IEEE Trans. Biomed. Eng.*, vol. 53, no. 12, pp. 2472–2478, 2006.
- [37] C. Dalla Man, R. A. Rizza, and C. Cobelli, “Meal simulation model of the glucose–insulin system,” *IEEE Trans. Biomed. Eng.*, vol. 54, no. 10, pp. 1740–1749, 2007.
- [38] A. Vaswani *et al.*, “Attention is all you need,” in *Advances in Neural Information Processing Systems*, vol. 30, 2017, pp. 5998–6008.
- [39] P. J. Huber, “Robust estimation of a location parameter,” *Ann. Math. Statist.*, vol. 35, no. 1, pp. 73–101, 1964.
- [40] R. Koenker and G. Bassett, Jr., “Regression quantiles,” *Econometrica*, vol. 46, no. 1, pp. 33–50, 1978.
- [41] S. Wang, Y. Teng, and P. Perdikaris, “Understanding and mitigating gradient flow pathologies in physics-informed neural networks,” *SIAM J. Sci. Comput.*, vol. 43, no. 5, pp. A3055–A3081, 2021.
- [42] A. Kendall, Y. Gal, and R. Cipolla, “Multi-task learning using uncertainty to weigh losses for scene geometry and semantics,” in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit.*, 2018, pp. 7482–7491.
- [43] W. L. Clarke *et al.*, “Evaluating clinical accuracy of systems for self-monitoring of blood glucose,” *Diabetes Care*, vol. 10, no. 5, pp. 622–628, 1987.
- [44] J. L. Parkes, S. L. Slatin, S. Pardo, and B. H. Ginsberg, “A new consensus error grid to evaluate the clinical significance of inaccuracies in the measurement of blood glucose,” *Diabetes Care*, vol. 23, no. 8, pp. 1143–1148, 2000.
- [45] A. Pfützner, D. C. Klonoff, S. Pardo, and J. L. Parkes, “Technical aspects of the Parkes error grid,” *J. Diabetes Sci. Technol.*, vol. 7, no. 5, pp. 1275–1281, 2013.
- [46] B. P. Kovatchev, D. J. Cox, L. A. Gonder-Frederick, and W. L. Clarke, “Symmetrization of the blood glucose measurement scale and its applications,” *Diabetes Care*, vol. 20, no. 11, pp. 1655–1658, 1997.
- [47] B. P. Kovatchev *et al.*, “Evaluation of a new measure of blood glucose variability in diabetes,” *Diabetes Care*, vol. 29, no. 11, pp. 2433–2438, 2006.
- [48] F. Wilcoxon, “Individual comparisons by ranking methods,” *Biometrics Bull.*, vol. 1, no. 6, pp. 80–83, 1945.
- [49] S. Holm, “A simple sequentially rejective multiple test procedure,” *Scand. J. Stat.*, vol. 6, no. 2, pp. 65–70, 1979.
- [50] P. E. Shrout and J. L. Fleiss, “Intraclass correlations: Uses in assessing rater reliability,” *Psychol. Bull.*, vol. 86, no. 2, pp. 420–428, 1979.
- [51] C. Spearman, “The proof and measurement of association between two things,” *Am. J. Psychol.*, vol. 15, no. 1, pp. 72–101, 1904.
- [52] T. Battelino *et al.*, “Clinical targets for continuous glucose monitoring data interpretation: Recommendations from the International Consensus on Time in Range,” *Diabetes Care*, vol. 42, no. 8, pp. 1593–1603, 2019.
- [53] S. Sivananthan, V. Naumova, C. Dalla Man, A. Facchinetti, E. Renard, C. Cobelli, and S. V. Pereverzyev, “Assessment of blood glucose predictors: The prediction-error grid analysis,” *Diabetes Technol. Ther.*, vol. 13, no. 8, pp. 787–796, 2011.
- [54] M. Sundararajan, A. Taly, and Q. Yan, “Axiomatic attribution for deep networks,” in *Proc. 34th Int. Conf. Mach. Learn.*, PMLR, vol. 70, 2017, pp. 3319–3328.
- [55] L. Breiman, “Random forests,” *Mach. Learn.*, vol. 45, pp. 5–32, 2001.
- [56] J. Martinsson, A. Schliep, B. Eliasson, C. Meijner, S. Persson, and O. Mogren, “Automatic blood glucose prediction with confidence using recurrent neural networks,” in *Proc. 3rd Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2148, pp. 64–68, 2018.
- [57] J. Xie and Q. Wang, “Benchmark machine learning approaches with classical time series approaches on the blood glucose level prediction challenge,” in *Proc. 3rd Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2148, pp. 97–102, 2018.
- [58] A. Bertachi, L. Biagi, I. Contreras, N. Luo, and J. Vehí, “Prediction of blood glucose levels and nocturnal hypoglycemia using physiological models and artificial neural networks,” in *Proc. 3rd Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2148, pp. 85–90, 2018.
- [59] K. Gu, R. Dang, and T. Prioleau, “Neural physiological model: A simple module for blood glucose prediction,” in *Proc. 42nd Annu. Int. Conf. IEEE Eng. Med. Biol. Soc.*, 2020, pp. 5476–5481.
- [60] J. Pavan *et al.*, “Personalized machine learning algorithm based on shallow network and error imputation module for an improved blood glucose prediction,” in *Proc. 5th Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2675, pp. 95–99, 2020.
- [61] R. Bevan and F. Coenen, “Experiments in non-personalized future blood glucose level prediction,” in *Proc. 5th Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2675, pp. 100–104, 2020.
- [62] H. Rubin-Falcone, I. Fox, and J. Wiens, “Deep residual time-series forecasting: Application to blood glucose prediction,” in *Proc. 5th Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2675, pp. 105–109, 2020.
- [63] J. Freiburghaus, A. Rizzotti-Kaddouri, and F. Albertetti, “A deep learning approach for blood glucose prediction and monitoring of type 1 diabetes patients,” in *Proc. 5th Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2675, pp. 131–135, 2020.
- [64] T. Yang *et al.*, “Multi-scale long short-term memory network with multi-lag structure for blood glucose prediction,” in *Proc. 5th Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2675, pp. 136–140, 2020.
- [65] T. Zhu, X. Yao, K. Li, P. Herrero, and P. Georgiou, “Blood glucose prediction for type 1 diabetes using generative adversarial networks,” in *Proc. 5th Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2675, pp. 90–94, 2020.
- [66] H. Nemat, H. Khadem, J. Elliott, and M. Benaissa, “Data fusion of activity and CGM for predicting blood glucose levels,” in *Proc. 5th Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2675, pp. 120–124, 2020.
- [67] H. Khadem, H. Nemat, J. Elliott, and M. Benaissa, “Multi-lag stacking for blood glucose level prediction,” in *Proc. 5th Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2675, pp. 146–150, 2020.
- [68] X. Sun, M. Rashid, M. Sevil, N. Hobbs, R. Brandt, M. R. Askari, A. Shahidehpour, and A. Cinar, “Prediction of blood glucose levels for people with type 1 diabetes using latent-variable-based model,” in *Proc. 5th Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2675, pp. 115–119, 2020.
- [69] M. Mayo and T. Koutny, “Neural multi-class classification approach to blood glucose level forecasting with prediction uncertainty visualisation,” in *Proc. 5th Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2675, pp. 80–84, 2020.
- [70] D. Joedicke, G. Kronberger, J. M. Colmenar, S. Winkler, J. M. Velasco, S. Contador, and J. I. Hidalgo, “Analysis of the performance of genetic programming on the Blood Glucose Level Prediction Challenge 2020,” in *Proc. 5th Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2675, pp. 141–145, 2020.
- [71] J. Daniels, P. Herrero, and P. Georgiou, “Personalised glucose prediction via deep multitask networks,” in *Proc. 5th Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2675, pp. 110–114, 2020.
- [72] N. Ma, Y. Zhao, S. Wen, T. Yang, R. Wu, R. Tao, X. Yu, and H. Li, “Online blood glucose prediction using autoregressive moving average model with residual compensation network,” in *Proc. 5th Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2675, pp. 151–155, 2020.

- [73] G. Cappon, L. Meneghetti, F. Prendin, J. Pavan, G. Sparacino, S. Del Favero, and A. Facchinetti, "A personalized and interpretable deep learning based approach to predict blood glucose concentration in type 1 diabetes," in *Proc. 5th Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2675, pp. 75–79, 2020.
- [74] A. R. Bhimoreddy, P. Sinha, B. Oluwalade, J. W. Gichoya, and S. Purkayastha, "Blood glucose level prediction as time-series modeling using sequence-to-sequence neural networks," in *Proc. 5th Int. Workshop Knowledge Discovery in Healthcare Data*, CEUR-WS, vol. 2675, pp. 125–130, 2020.
- [75] C. Piao, T. Zhu, S. E. Baldeweg, P. Taylor, P. Georgiou, J. Sun, J. Wang, and K. Li, "GARNN: An interpretable graph attentive recurrent neural network for predicting blood glucose levels via multivariate time series," *Neural Netw.*, vol. 185, Art. no. 107229, 2025, doi: 10.1016/j.neunet.2025.107229.